

Particle Swarm Optimization (PSO)

1. Introduction

Particle Swarm Optimization (PSO) is a **swarm intelligence** method inspired by collective behaviors observed in nature, such as flocks of birds or schools of fish. It is widely applied in various optimization problems, including **course scheduling, function optimization, and machine learning hyperparameter tuning**. The algorithm operates by maintaining a population (or swarm) of candidate solutions, known as particles, that iteratively adjust their positions in search of an optimal solution.

This study implements two key variants of PSO:

1. **Global Best PSO (Gbest PSO):**

- In this approach, each particle's movement is influenced by **the globally best solution** found within the swarm.
- This structure allows **rapid convergence**, as particles are collectively guided toward the most promising solution.
- However, it may increase the risk of **premature convergence**, where particles become trapped in suboptimal solutions due to lack of diversity.

2. **Local Best PSO (Lbest PSO):**

- Instead of relying on a single global best, **each particle interacts only with its local neighbors**, forming dynamic neighborhoods.
- This decentralized learning promotes **diversity** in the search process, reducing the likelihood of stagnation in suboptimal solutions.
- The localized structure enables **gradual refinement** of solutions, improving robustness while maintaining adaptability.

2. Algorithm Design

2.1. Particle Representation

Each particle in the swarm represents a **possible scheduling solution**, structured as:

- **Course ID, Instructor ID, Room Assignment, Timeslot**
- This ensures each candidate solution is a **valid representation** within the scheduling constraints.

2.2. Swarm Initialization

- The swarm consists of **S particles**, each initialized with **random room and timeslot assignments**.

- Each particle maintains a velocity vector, which dictates **how it moves** in the scheduling space.

2.3. Neighborhood Formation

Unlike Global Best PSO (Gbest PSO), where all particles are influenced by the **global best**, this implementation:

- **Divides the swarm into N neighborhoods** (groups of particles).
- **Each particle selects k nearest neighbors** based on Manhattan distance (i.e., similarity of room and timeslot scheduling).
- The particle follows **the best neighbor**, rather than a global best solution.

2.4. Velocity Update Mechanism

Each particle updates its velocity using the equation:

$$v_i(t + 1) = \omega v_i(t) + c_1 r_1 (pbest - x_i) + c_2 r_2 (nbest - x_i) v_i(t + 1)$$

Where:

- $v_i(t + 1)$ is the **updated velocity** of particle i at iteration (t + 1).
- $v_i(t)$ is the **current velocity** of particle i.
- x_i is the **current position** of particle i.
- ω (inertia weight) controls how much of the previous velocity is retained.
- c_1 (cognitive coefficient) influences the particle's movement toward its own best solution.
- c_2 (social coefficient) guides the particle toward the **best solution among its k-nearest neighbors**.
- r_1, r_2 are random values between [0,1] to introduce stochastic behavior.
- $pbest$ is the **personal best position** of the particle.
- $nbest$ is the **best position found by its local neighbors**.

2.5. Position Update

Once velocity is updated, the new **position** of the particle is computed as:

$$x_i(t + 1) = x_i(t) + v_i(t + 1)$$

This ensures the particle moves based on the learned velocity while exploring new solutions.

2.6. Adaptive Neighborhoods

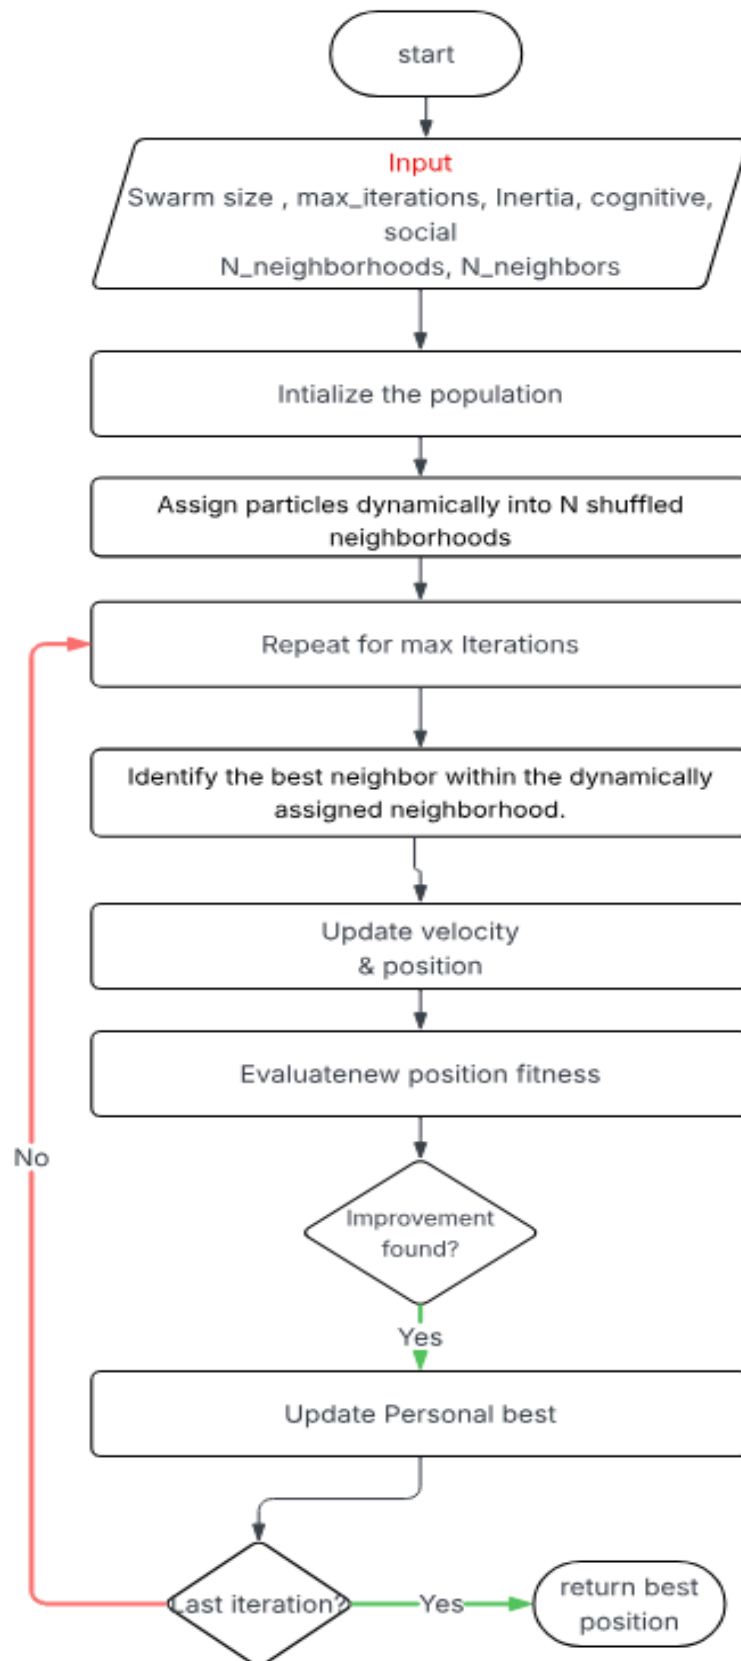
Neighborhoods are shuffled dynamically after initialization to **prevent stagnation** in local optima.

2.7. Algorithm parameters setup

Subset of Data (20%)	
PSO Swarm size	100 Particles
PSO Iterations Number	1000 Iterations
Inertial weight	0.5
Cognitive Learning Rate	1
Social Learning Rate	1
Number of Neighborhoods	10
Neighbors per particle	7

All Data	
PSO Swarm size	250 Particles
PSO Iterations Number	1000 Iterations
Inertial weight	0.5
Cognitive Learning Rate	2
Social Learning Rate	2
Number of Neighborhoods	10
Neighbors per particle	7

3) Lbest PSO Approach Flowchart



4) Lbest PSO Approach Pseudocode

Algorithm: Local Best Particle Swarm Optimization (lbest PSO)

Input:

- **S** = swarm size (number of particles)
- **max_iterations** = number of optimization cycles
- **N** = number of neighborhoods (subdivisions of swarm)
- **k** = number of neighbors per neighborhood (local best selection)
- **inertia, cognitive, social** = PSO hyperparameters controlling movement behavior

Initialization:

1. **Create a swarm of S particles**, each representing a potential scheduling solution.
2. Assign each particle a **random position** (course assignments to rooms and timeslots).
3. Assign each particle a **random velocity**, affecting room and timeslot adjustments.
4. Set each particle's **best personal position** to its initial state.
5. Compute **fitness** for each particle and update the best-known global solution.
6. Assign particles dynamically into **N shuffled neighborhoods** to ensure varied local interactions.

Optimization Process (Repeat for max_iterations):

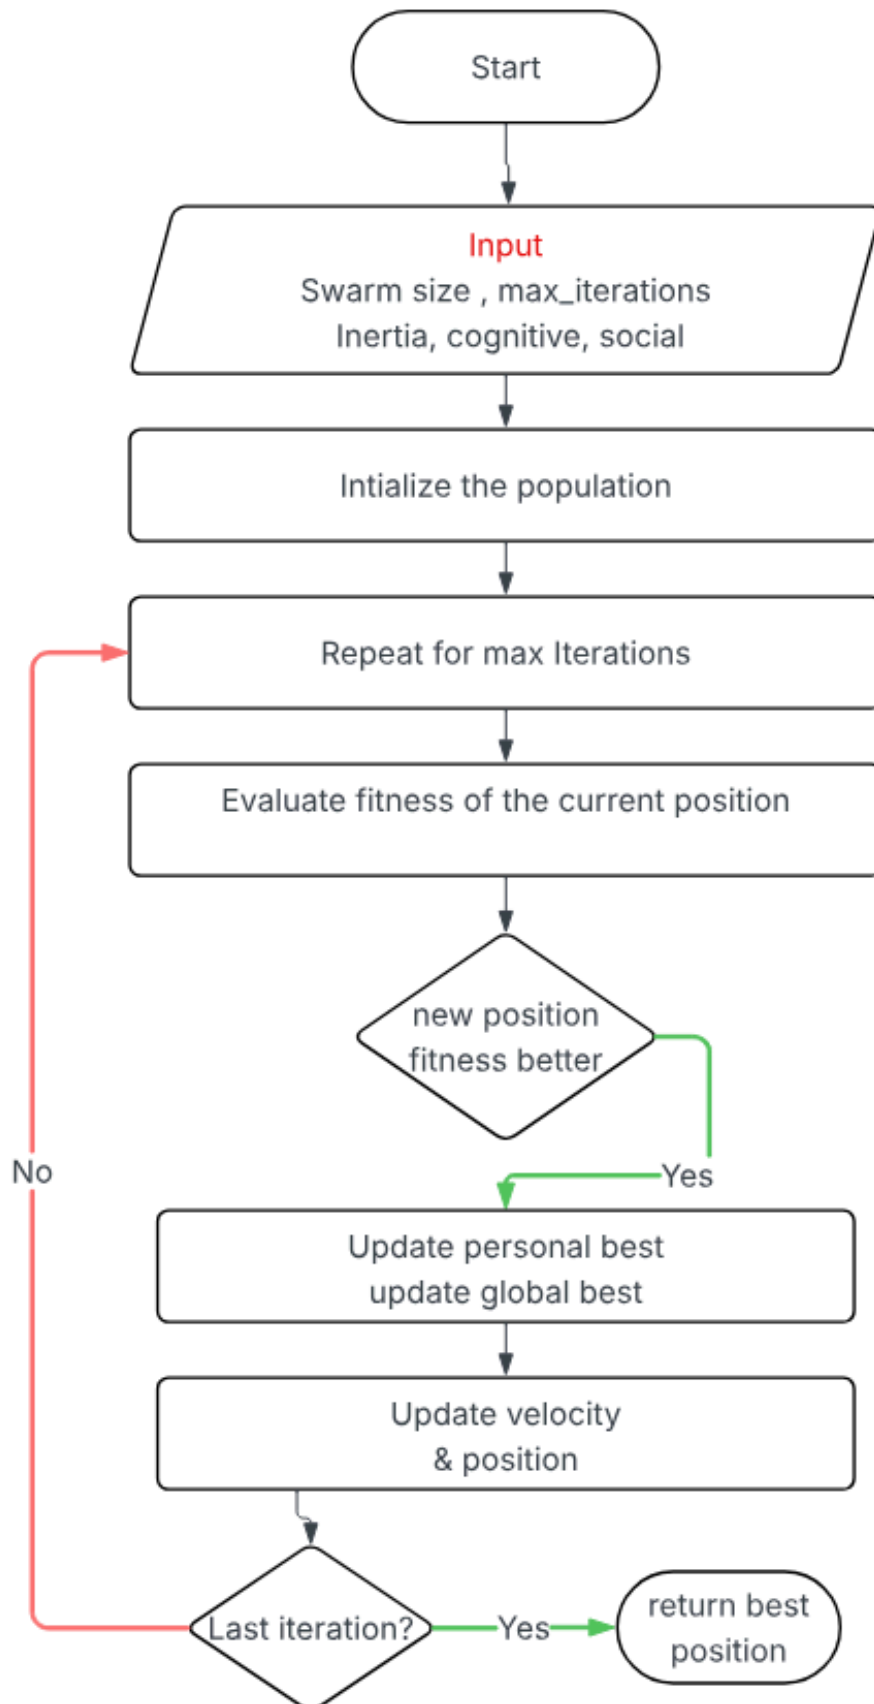
For each particle in the swarm:

1. **Identify the best neighbor** within the dynamically assigned neighborhood.
2. **Update velocity** using:
 - **Inertia component** (retains some influence from previous velocity).
 - **Cognitive component** (adjusts towards the particle's personal best).
 - **Social component** (adjusts towards the best-performing neighbor).
3. **Update position** based on new velocity, ensuring valid constraints (room availability, timeslot restrictions).
4. **Evaluate new position's fitness** and update personal best if an improvement is found.
5. Track **global best solution** across all particles while neighborhoods continue evolving.

Return:

- The **best scheduling solution** found during optimization.

5) Gbest PSO Approach Flowchart



6) Gbest PSO Approach Pseudocode

Algorithm: Global Best Particle Swarm Optimization (gbest PSO)

Input:

- **S** = swarm size (number of particles)
- **max_iterations** = number of optimization cycles
- **inertia, cognitive, social** = PSO hyperparameters controlling movement behavior

Initialization:

1. **Create a swarm of S particles**, each representing a potential scheduling solution.
2. Assign each particle a **random position** (course assignments to rooms and timeslots).
3. Assign each particle a **random velocity**, affecting room and timeslot adjustments.
4. Set each particle's **best personal position** to its initial state.
5. Compute **fitness** for each particle using a predefined fitness function.
6. Identify the best-performing particle and set it as the **global best solution**.

Optimization Process (Repeat for max_iterations):

For each particle in the swarm:

1. **Evaluate fitness** to determine how good the current schedule is.
2. **Update personal best** if the new position provides a better fitness score.
3. **Update global best** if the particle's personal best is better than the current global best.
4. **Update velocity** using:
 - **Inertia component** (retains some influence from previous velocity).
 - **Cognitive component** (adjusts towards the particle's personal best).
 - **Social component** (adjusts towards the global best solution).
5. **Update position** based on new velocity, ensuring constraints (room availability, valid timeslots).
6. Repeat the process across all particles until convergence or reaching **max_iterations**.

Return:

- The **best scheduling solution** found during optimization.

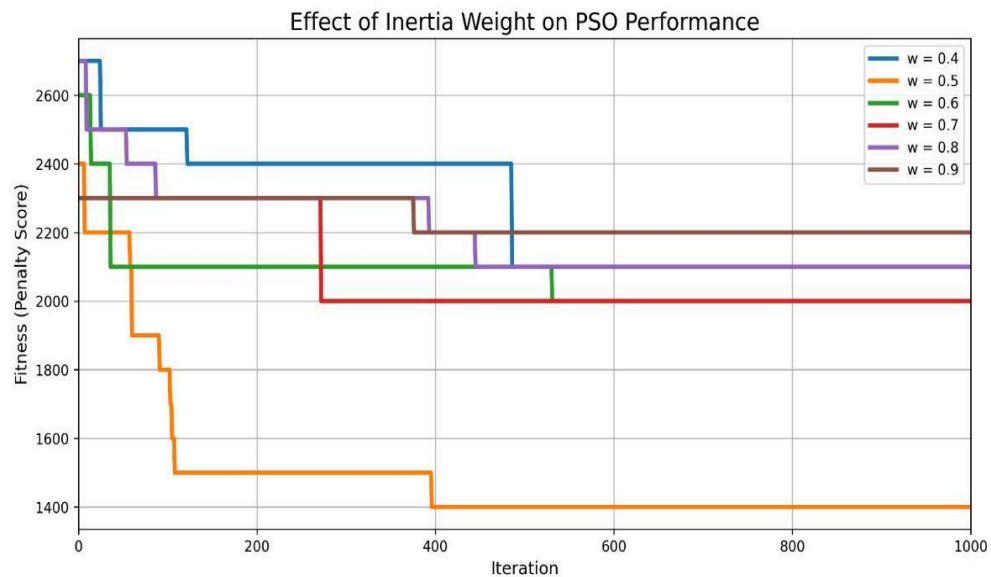
7. Experimental Setup & Results

The results you obtained are based on only 20% of the available data. This limitation occurred due to constraints in computational power and the excessively long processing time required to analyze the full dataset. As a result, the findings might not fully represent the entire dataset and could be impacted by the reduced sample size.

3.1. Effect of Inertia weight on PSO Performance

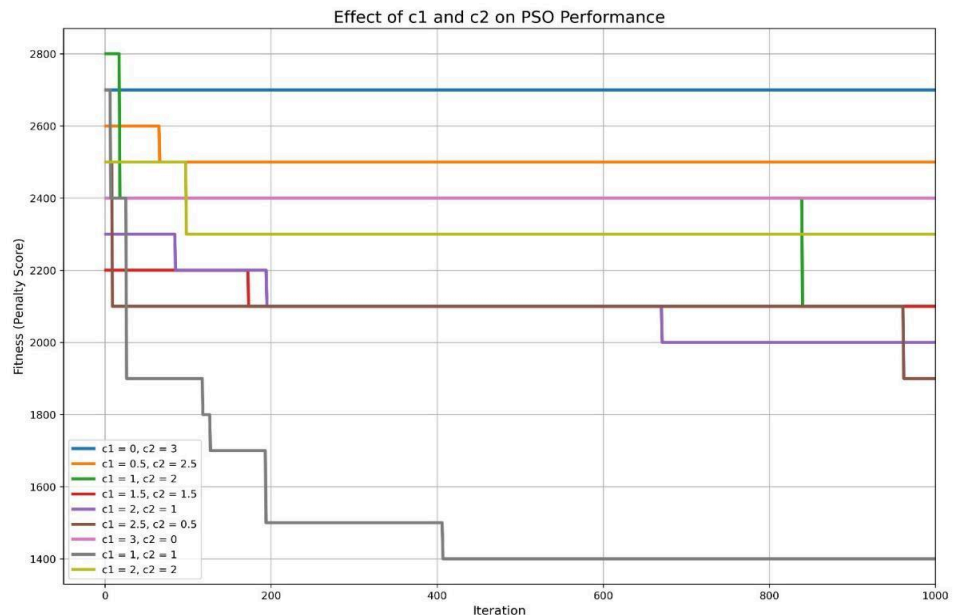
The inertia weight (ω) plays a significant role in determining the performance of Particle Swarm Optimization (PSO) fitness, though its behavior is not entirely fixed or predictable. However, a general trend emerges: fitness improvements tend to increase as ω decreases.

In our experiment, most ω values exhibited early stabilization or convergence, typically reaching that state around the halfway mark of the maximum iterations. Based on the results, setting ω around **0.5** appears to offer an optimal balance. This value demonstrated the best overall performance—both in terms of final fitness score and the level of improvement observed throughout the iterations.



3.2. Effect of C1 and C2 combinations on PSO Performance

The impact of varying the coefficients c_1 and c_2 on fitness performance in Particle Swarm Optimization (PSO) appears to be relatively minor, as most parameter combinations exhibit similar effects. However, assigning a value of **3** to either c_1 or c_2 tends to degrade performance, with fitness values remaining constant across iterations, indicating stagnation.

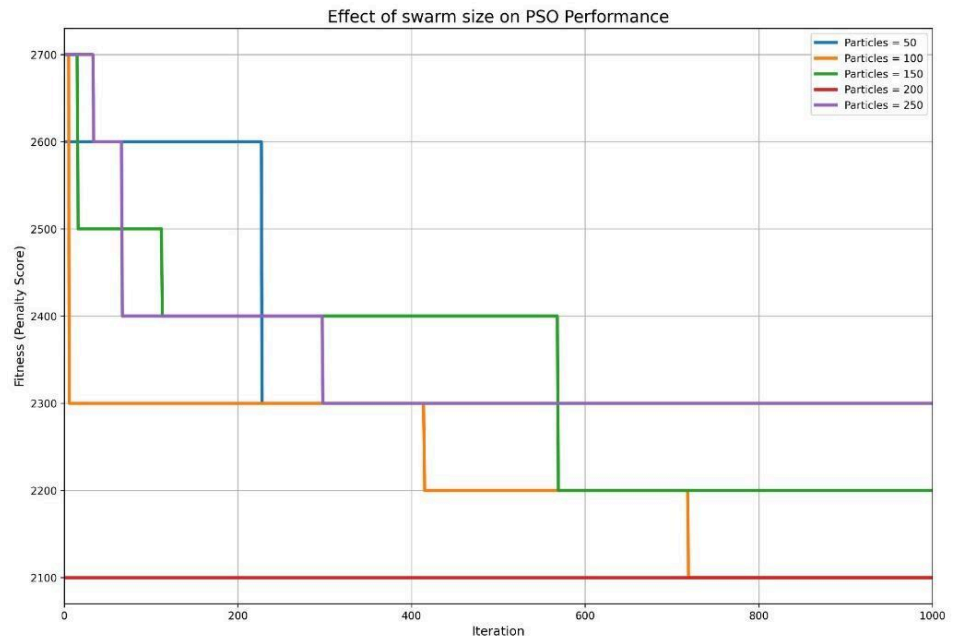


Conversely, parameter combinations where **one of the values is set to 1** result in noticeably improved fitness performance. These configurations demonstrate both higher fitness improvements and better final scores across generations. Based on the experimental results, setting $c_1 = 1$ and $c_2 = 1$ emerges as the most effective configuration, yielding the best overall performance in terms of both convergence behavior and final fitness values.

3.3. Effect of Swarm Size on PSO Performance

Among the different swarm sizes tested, **100 particles consistently provided the best optimization**, demonstrating strong initial fitness and efficient convergence to a competitive final score. The **150 and 250 particle configurations** also showed steady improvement throughout the iterations, maintaining consistent progress before stabilizing.

On the other hand, the **50-particle swarm reached convergence early**, with only minimal improvement, leading to a higher penalty score. The **200-particle swarm exhibited the weakest performance**, showing limited optimization and stagnation at a higher fitness value.



For this particular experiment, a **swarm size of 100 particles** emerged as the most effective configuration, offering the best balance between performance and convergence behavior. **Swarm sizes of 150 and 250 also performed reasonably well**, whereas **50 and 200 particles showed weaker results due to early convergence or stagnation**.


3.4. Effect of Number of neighborhood on PSO Performance

All neighborhood configurations began with high fitness scores and exhibited a consistent decrease in penalty throughout the generations. However, **most configurations experienced early convergence**, suggesting limited exploration in the later stages of the optimization process.

Neighborhood Size = 7: Achieved the **fastest convergence** while maintaining the **best fitness score**, demonstrating a strong balance between **exploration** (global search) and **exploitation** (local refinement).

Neighborhood Size = 10: Showed the **most substantial fitness improvement before convergence**, suggesting **better initial exploration** but **weaker exploitation** toward the final solution.

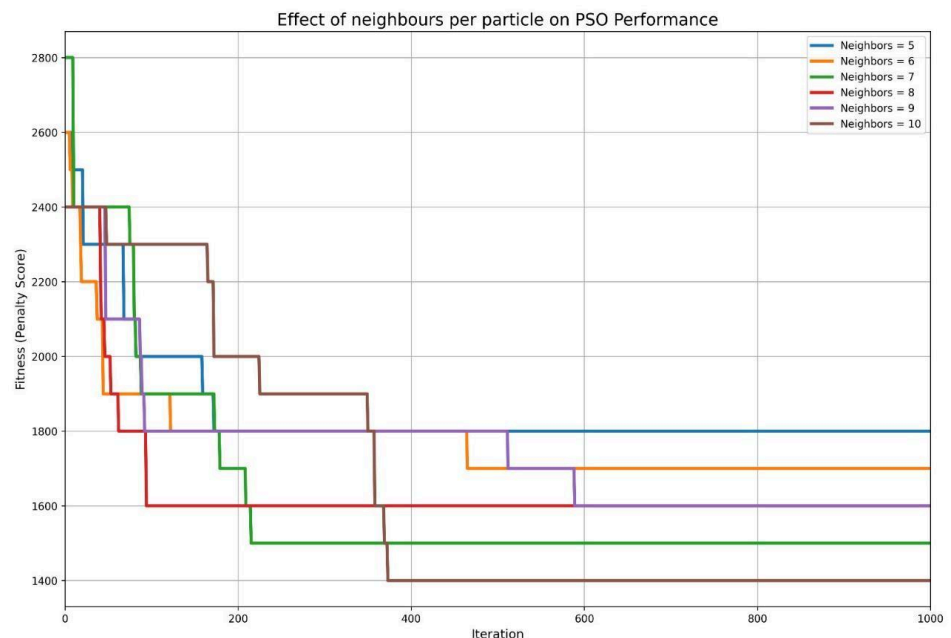
Neighborhood Sizes = 5, 6, 8, and 9: Converged relatively early without reaching particularly low fitness values. These configurations appeared to stagnate after a short period of improvement, indicating **either weak exploration or premature convergence**.



Among the tested configurations, a **neighborhood size of 7** provided the best trade-off between exploration and exploitation, resulting in **efficient optimization**. Meanwhile, a **neighborhood size of 10** exhibited strong initial improvements but weaker refinement in later stages. Other configurations displayed premature convergence, limiting their overall effectiveness.

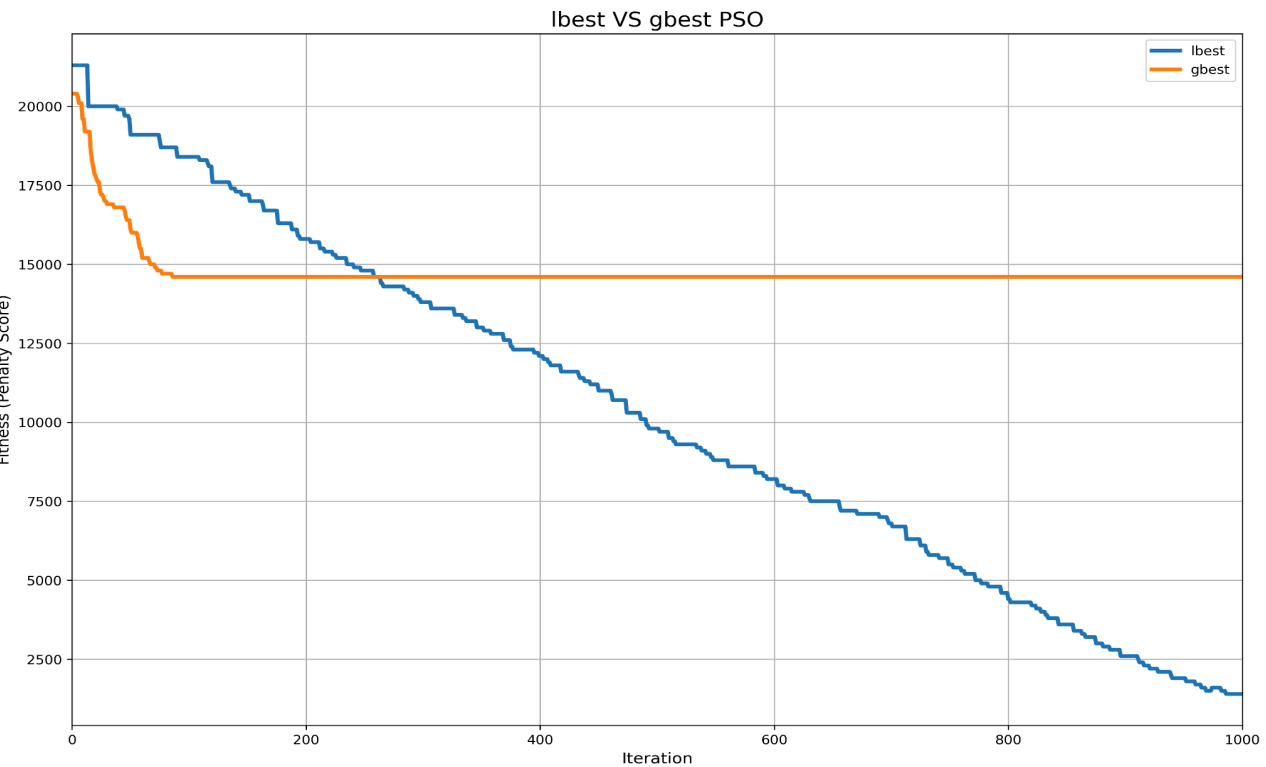
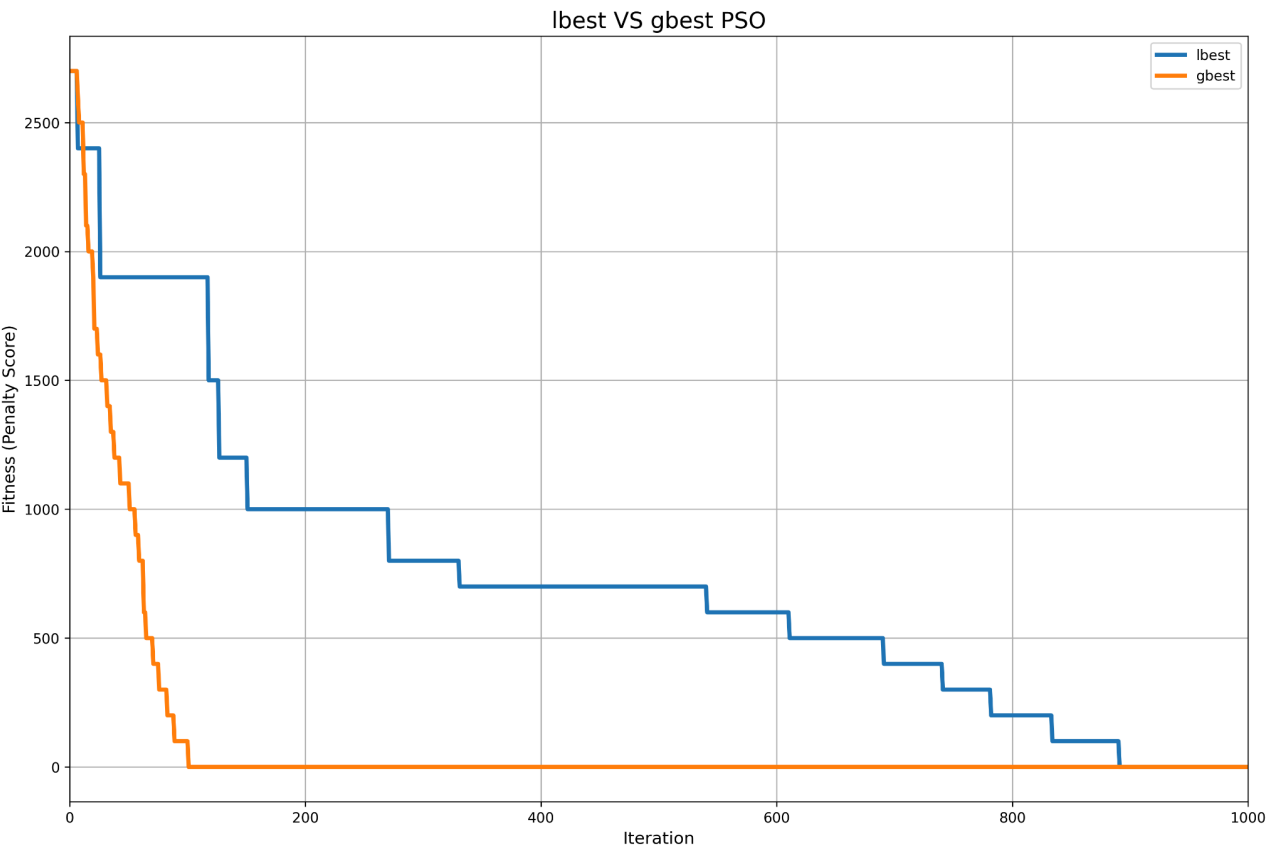
3.5. Effect of neighbors per particle on PSO Performance

The number of neighbors per particle influences the balance between exploration and exploitation in PSO. Configurations with **5, 6, 8, and 9 neighbors** exhibited early convergence but stagnated at higher fitness values, suggesting limited exploration. **A neighbors size of 7** achieved the fastest convergence and the lowest fitness score, demonstrating an optimal balance between search diversity and refinement. In contrast, **10 neighbors showed prolonged improvement before settling**, indicating strong initial exploration but weaker exploitation.



A **neighborhood size of 7** provided the most efficient optimization, offering a trade-off between global search and local refinement. Smaller and larger neighborhood configurations either stagnated prematurely or took longer to converge, affecting overall performance.

3.5. Gbest VS Lbest approaches



Particle Swarm Optimization (PSO) is known for its ability to balance exploration and exploitation in search spaces. This study compares **Global Best PSO (Gbest PSO)** and **Local Best PSO (Lbest PSO)** based on their optimization behavior, analyzing how each approach influences convergence, search diversity, and final fitness values.

8. Observations from the Results

8.1. Global Best PSO (Gbest PSO)

- **Initial Performance:** Gbest PSO starts with **high fitness values**, indicating strong global guidance toward optimal solutions.
- **Convergence Behavior:** The rapid improvement in early iterations suggests **efficient exploitation**, allowing particles to quickly refine solutions.
- **Stagnation Risks:** While Gbest PSO achieves significant gains early, its reliance on a global best leader can lead to **premature convergence**, limiting exploration diversity.
- **Final Fitness Values:** The final optimization reaches competitive fitness levels, but the steady decline suggests that the swarm loses some adaptability after early improvements.

8.2. Local Best PSO (Lbest PSO)

- **Initial Performance:** Lbest PSO exhibits **gradual fitness improvement** over a longer period, emphasizing **localized learning** rather than immediate global influence.
- **Convergence Behavior:** The extended improvement phases suggest **better exploration**, allowing particles to search effectively before settling.
- **Adaptability and Diversity:** Unlike Gbest PSO, Lbest PSO maintains **higher variance in fitness values throughout iterations**, promoting flexibility in solution refinement.
- **Final Fitness Values:** The final solutions demonstrate a **smoother convergence curve**, indicating stronger adaptability across diverse search areas.

8.3 Behavioral Analysis

- **Exploration vs. Exploitation:** Gbest PSO prioritizes exploitation early, leading to faster but sometimes restrictive convergence. Lbest PSO, in contrast, ensures **wider search coverage** due to decentralized learning, reducing the risk of stagnation.
- **Search Stability:** Lbest PSO maintains adaptability, benefiting complex search spaces where premature convergence is undesirable.
- **Optimization Efficiency:** Gbest PSO is **ideal for problems requiring rapid convergence**, while Lbest PSO is **better suited for scenarios needing a broader search before refinement**.